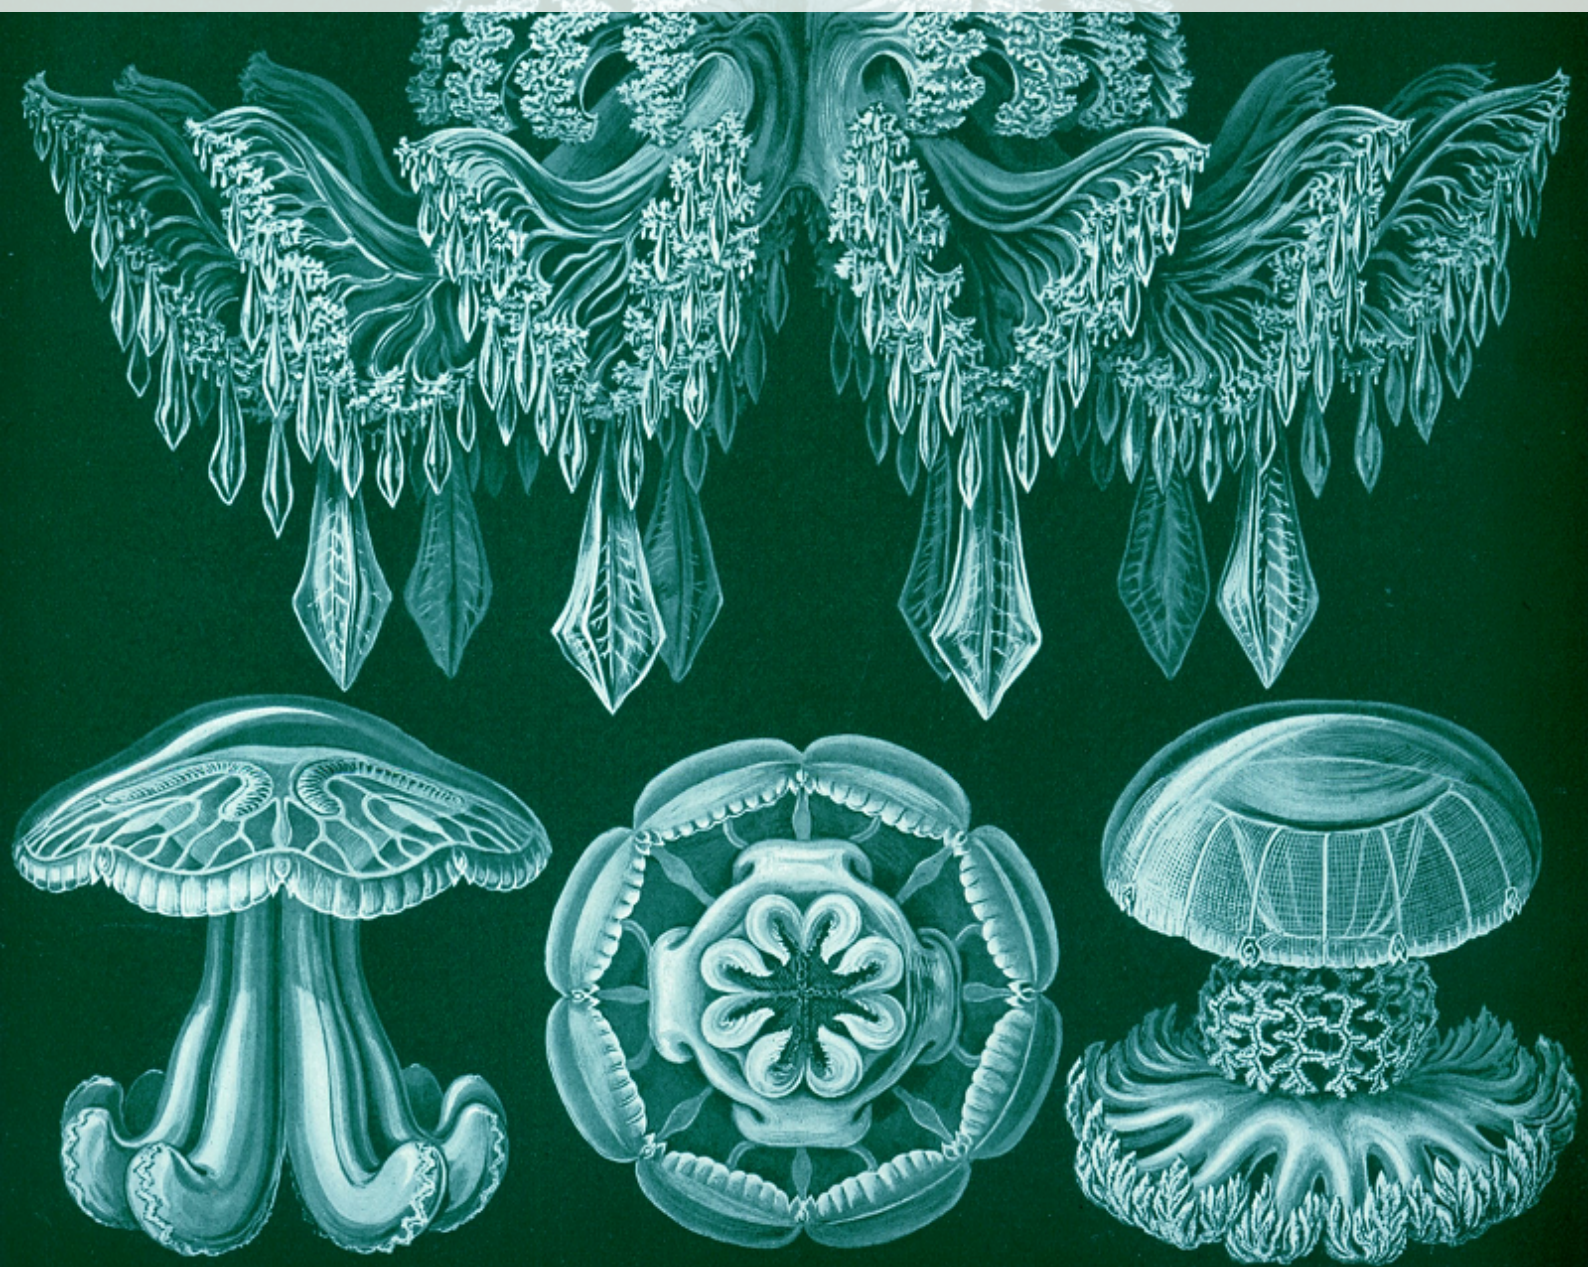




Oracle Trace Files Explained

Attempting to Document the Internals of an Oracle Trace File

Norman Dunbar



Copyright ©2017,2018 Norman Dunbar

PUBLISHED BY MEMYSELF EYE PUBLISHING ;-)

The latest version of the eBook can be downloaded from [the book's github repository](#).

Licensed under the Creative Commons Attribution-NonCommercial 3.0 Unported License (the “License”). You may not use this file except in compliance with the License. You may obtain a copy of the License at <http://creativecommons.org/licenses/by-nc/3.0>. Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an “AS IS” BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

First printing, July 2017

This pdf document was created on 21/6/2018 at 10:12:29.



Contents

| I | Oracle Trace Files | |
|-------|---|----|
| 1 | Introduction | 3 |
| 2 | Trace File Gotchas! | 5 |
| 2.1 | Things to Beware Of | 5 |
| 2.1.1 | When Are Lines Written to the Trace File? | 5 |
| 2.1.2 | Tim Values | 5 |
| 2.1.3 | Timestamps | 6 |
| 2.1.4 | Cursor Depth | 6 |
| 2.1.5 | Cursor IDs Get Reused | 6 |
| 2.1.6 | What About Child Cursors? | 7 |
| 3 | Oracle Trace Files Explained | 9 |
| 3.1 | Trace File Sections | 9 |
| 3.1.1 | The Header | 9 |
| 3.1.2 | Trace Details | 10 |
| 3.2 | Trace File Details | 13 |

| | | |
|----------|---------------------------------|-----------|
| 4 | Trace File Entries | 15 |
| 4.1 | PARSING IN CURSOR | 15 |
| 4.2 | PARSE | 17 |
| 4.3 | PARSE ERROR | 19 |
| 4.4 | EXEC | 20 |
| 4.5 | FETCH | 22 |
| 4.6 | WAIT | 23 |
| 4.7 | ERROR | 25 |
| 4.8 | STAT | 26 |
| 4.9 | CLOSE | 28 |
| 4.10 | XCTEND | 29 |
| 4.11 | BINDS | 30 |
| 4.11.1 | Examples | 31 |
| 4.12 | UNMAP | 34 |
| 4.13 | SORT UNMAP | 35 |

II

Appendices

| | | |
|----------|--|-----------|
| A | Oracle Data Types | 39 |
| B | Oracle Command Codes | 43 |
| C | Oracle Characterset Codes | 47 |
| D | How this Book Evolved | 51 |
| D.1 | Why this Book? | 51 |
| D.2 | Creating the Book | 51 |
| E | TraceAdjust Utility | 53 |
| F | TraceMiner2 Utility | 55 |



List of Tables

| | | |
|------|--|----|
| 4.1 | Parsing in Cursor - Fields. . . <i>continues on next page</i> | 15 |
| 4.1 | Parsing in Cursor - Fields | 16 |
| 4.2 | Parse - Fields | 17 |
| 4.3 | Parse Error - Fields | 19 |
| 4.4 | Exec - Fields | 20 |
| 4.5 | fetch - Fields | 22 |
| 4.6 | Wait - Fields for Oracle 9i | 23 |
| 4.7 | Parse - Fields for Oracle 10g Onwards | 24 |
| 4.8 | Error - Fields | 25 |
| 4.9 | Stat - Fields for Oracle 9i | 26 |
| 4.10 | Stat - Fields for Oracle 10g Onwards | 27 |
| 4.11 | Close - Fields | 28 |
| 4.12 | Xctend - Fields | 29 |
| 4.13 | Binds - Fields for Oracle 9i | 30 |
| 4.14 | Binds - Fields for Oracle 10g Onwards. . . <i>continues on next page</i> | 30 |
| 4.14 | Binds - Fields for Oracle 10g Onwards | 31 |
| 4.15 | Parse - Fields | 34 |
| 4.16 | Parse - Fields | 35 |
| A.1 | Bind Variable data Types. . . <i>continues on next page</i> | 39 |
| A.1 | Bind Variable data Types | 40 |
| B.1 | Oracle Command Codes. . . <i>continues on next page</i> | 43 |
| B.1 | Oracle Command Codes. . . <i>continues on next page</i> | 44 |
| B.1 | Oracle Command Codes | 45 |
| C.1 | Oracle Character Set Codes. . . <i>continues on next page</i> | 47 |

| | | |
|-----|---|----|
| C.1 | Oracle Character Set Codes... <i>continues on next page</i> | 48 |
| C.1 | Oracle Character Set Codes | 49 |



Listings

| | | |
|------|---|----|
| 2.1 | Counting Child Cursors in a Trace File | 7 |
| 2.2 | Extracting Child Cursors from a Trace File | 7 |
| 3.1 | Oracle 11g Trace File Header | 9 |
| 3.2 | Oracle 11g Trace File Header - With Waits etc | 10 |
| 3.3 | Oracle 11g Problem Trace - grep output | 11 |
| 3.4 | Time Stamp Line | 11 |
| 4.1 | Parsing In Cursor Line | 15 |
| 4.2 | Parsing In Cursor Line - Proxy User | 16 |
| 4.3 | Parse Line | 17 |
| 4.4 | Parse Error Line | 19 |
| 4.5 | Exec Line | 20 |
| 4.6 | PL/SQL Exec Example | 20 |
| 4.7 | Fetch Line | 22 |
| 4.8 | Wait Line - Oracle 9i | 23 |
| 4.9 | Extracting Event Names for Oracle 9i | 23 |
| 4.10 | Wait Line - Oracle 10g Onwards | 23 |
| 4.11 | Error Line | 25 |
| 4.12 | Stat Line | 26 |
| 4.13 | Close Line | 28 |
| 4.14 | Commit Statement with Xctend Line | 29 |

| | | |
|------|--|----|
| 4.15 | Example of Recursive SQL Statement | 31 |
| 4.16 | Binds Lines | 31 |
| 4.17 | Bind Example - VARCHAR2 with WE8ISO8859P1 Characterset | 32 |
| 4.18 | Bind Example - VARCHAR2 with ALUTF16 Characterset | 32 |
| 4.19 | Bind Example - REF_CURSOR | 32 |
| A.1 | Bind Example - ROWID With oacdtty=11 | 40 |
| B.1 | SQL Query to List Oracle Command Codes | 45 |
| C.1 | SQL Query to list Character Set Codes and Names | 47 |
| C.2 | SQL Query to Convert a csi Code to a Charcter Set Name | 49 |
| C.3 | SQL Query to Convert a Charcter Set Name to a csi Code | 49 |
| E.1 | TraceAdjust Example - Before Processing | 53 |
| E.2 | TraceAdjust Example - After Processing | 54 |

Oracle Trace Files

| | | |
|----------|---|-----------|
| 1 | Introduction | 3 |
| 2 | Trace File Gotchas! | 5 |
| 2.1 | Things to Beware Of | |
| 3 | Oracle Trace Files Explained | 9 |
| 3.1 | Trace File Sections | |
| 3.2 | Trace File Details | |
| 4 | Trace File Entries | 15 |
| 4.1 | PARSING IN CURSOR | |
| 4.2 | PARSE | |
| 4.3 | PARSE ERROR | |
| 4.4 | EXEC | |
| 4.5 | FETCH | |
| 4.6 | WAIT | |
| 4.7 | ERROR | |
| 4.8 | STAT | |
| 4.9 | CLOSE | |
| 4.10 | XCTEND | |
| 4.11 | BINDS | |
| 4.12 | UNMAP | |
| 4.13 | SORT UNMAP | |



1. Introduction

Oracle trace files are not greatly documented. This document is an attempt to do so. It is *not* official in any way and is based on a good few years of reading these files to help diagnose various database problems.

A trace file is really the best way to delve into what Oracle is doing, or to discover why something is taking so long - it shows you exactly what happened during the period that the session was being traced.

Even better, when you extract an Explain Plan from a trace file, it is showing you exactly how Oracle retrieved the data from the tables, and exactly where the time was spent in doing so. Running an Explain Plan for ... statement in *SQL*Plus*, *Toad* etc tells you what Oracle *might* do. The two are not always the same.

The trace files used (and abused) in this eBook were all self contained, and extracted from a dedicated session on the database. If you are using shared sessions, then you have the unfortunate problem of having bits of your session trace spread over any number of different trace files. Not fun.

From Oracle 10g onwards, a utility has been supplied, named *trcsess*. This allows you to collect together all the separate files and merge them into one, ready for analysis. Even better, it apparently can also merge Oracle 9i shared server session trace files. Bonus!

Trace File Extracts

In the remainder of this eBook, extracts from trace files will be shown as follows:

```
WAIT #3220341128: nam='db file sequential read' ela= 1023 file#=3 block  
=>#=12 blocks=1 obj#=-1 tim=3520817183625
```


You will notice that long lines from the trace file are wrapped in the example above. However, all such wrapped lines are indicated by a \implies at the start of the continuation line(s). Sometimes though, the break in the lines can be at an awkward position. Sadly, I can't help this, it's all automagically done and I have no input as to how, exactly, \LaTeX decides.

Code Listings

SQL scripts (or any other program listings) will be listed as per the example below. The only difference being line numbers, in case I have to reference a particular point, and syntax highlighting.

```
1 select name, parameter1, parameter2, parameter3
2 from v$event_name
3 where name = 'db file sequential read';
```

Irregular Updates

This eBook is a work in progress, I am always learning, or trying to. As I update it, the latest version will always appear on GitHub as detailed on the Copyright page above. Also on that page, you will be able to see the date and time that the version of the eBook you are reading, was generated. That's done automagically too, so I don't need to remember to update it - thankfully!

Comments

Seen anything blatantly wrong? Something you think needs a better explanation? Too many blatant plugs? You can either create an issue on GitHub, or, [email me](#) with details and I'll do my best to get it fixed.



2. Trace File Gotchas!

Trace files can be full of pitfalls and traps for the unwary. Best take care!

2.1 Things to Beware Of

There are a few things to watch out for in trace files, some of these are mentioned frequently in the text below, but they are also gathered here together, in one place, for your convenience.¹

2.1.1 When Are Lines Written to the Trace File?

It cannot be said often enough, *a line in the trace file is written when that particular action has completed* and not before. So, if you see a line defining a PARSE statement in the trace file, the PARSE has completed at that point. The `tim` value gives you the (relative) time that the action completed.

2.1.2 Tim Values

On 11g `tim` values are in microseconds, or, millionths of a second, and are an offset from some epoch, which itself depends on the Operating System in use.

You can get a rough idea of how long a statement took by subtracting the previous `tim` from the current `tim` - the answer is in microseconds, so count back 6 places from the right end of the `tim` value, and shove in a decimal point.

Alternatively, use my `TraceAdjust` utility to convert `tim` values into date and times accurate to the microsecond. (Plug! See Appendix E on page 53 for details.)

¹Ok, for *my* convenience!

2.1.3 Timestamps

Not all actions in the trace file may be *accurately* timed. This can lead to the various `tim` values getting slightly out of step. To attempt to alleviate this, Oracle will write timestamp lines to the trace file. These are flagged by three asterisks (***) followed by the date and time, accurate to microsecond level, possibly in the format `yyyy-mm-dd hh24:mi:ss.u` but this *might* depend on locality etc. The `tim` value immediately following these timestamp lines are adjusted to match the timestamp that preceded it.

You can see an example of a timestamp line in the header for the trace file. An example is in Section 3.1.1 on Page 9.

2.1.4 Cursor Depth

Cursors are parsed, executed etc at a given depth. This can be seen in the `dep=n` details in the appropriate lines in the trace file.

If the listed depth is zero, then this is the top level SQL in the trace file, and usually corresponds to your own statements. Statements running in a PL/SQL procedure, function etc, however, *may* be listed at a different depth with the call to the procedure (or whatever) being run at depth zero.

As mentioned above, lines are written to the trace file when an action completes, so you will see (at least in most trace files) the PARSE, EXEC and FETCH and possibly STAT and/or CLOSE for the recursive SQL statements *after* the PARSE but *before* the EXEC for your statement. Simply because of the need to execute all those recursive statements in order to facilitate the EXEC of yours.

If you see a statement being actioned at `dep=n`, where $n \neq 0$, then it is being actioned recursively for another statement at `dep=n - 1` which *follows later* in the trace file.

2.1.5 Cursor IDs Get Reused

When you see something like `PARSE #1234567890` then the digits after the # are the cursor ID. In the old days, prior to Oracle 9i, cursor IDs were monotonically increasing and were never reused in the same trace file. Once cursor #1 had been closed, you wouldn't see it again.

This is no longer true and a cursor ID *can* be reused once closed. There are two (known to me) variations:

Same SQL Statement Reused by Same CursorID

Some applications write SQL that is parsed far too frequently and may even be parsed each and every time it is executed by the application code. If this is the case, the trace files *may* show the following steps

- CLOSE of the current cursor.
- *Possibly* a PARSE for the new statement. This will not be present if the previous CLOSE resulted in the cursor being cached.
- And so on.

If the same statement is reparsed and nothing else has been parsed since the CLOSE, then it is possible that the expected `PARSING IN CURSOR` line, showing the SQL text, may be missing from the trace file.

If the same statement has been parsed three or more times already, Oracle may cache the cursor rather than hard CLOSEing it as this is a performance improvement if the SQL has to be PARSED

and EXECuted again. In this case, there will be no PARSE line in the trace file for the statement - simply because it was not actually PARSEd, merely retrieved from the cache already PARSEd.

Different SQL using Same Cursor ID

If a statement's cursor is closed and another *different* statement is parsed, it *can* obtain the same Cursor ID as the recently closed cursor. In this case, you will see the following:

- CLOSE of the current cursor.
- PARSING IN CURSOR and the SQL text of the new statement.
- PARSE for the new statement.
- And so on.

The moral to this tale is simple, you cannot be sure that all the EXECs for a given Cursor ID actually mean that the exact same SQL statement has been executed - you could be seeing lots of different SQL statements simply reusing the same Cursor ID.

2.1.6 What About Child Cursors?

Following on from the above, the same SQL statement *can* (and does) have the ability to be executed under any number of *different* Cursor IDs! How is this possible?

If your statement uses bind variables, then given a recent version of Oracle, something called *Bind Variable Peeking* may be in use to ensure that a statement has the most efficient plan of execution regardless of the bind variable values. In the good old days, whatever bind value got in first determined the execution plan for every other EXEC of that statement, even if it made the execution plan completely hopeless for most values that followed.

I think - because I have not checked - that this means that certain statements and bind combinations will cause a number of separate child cursors, and thus, different Cursor IDs in the trace file, for the same statement. A recent trace file I examined had 15 different Cursor IDs for the same statement.

I knew they were all the same as the PARSING IN CURSOR line has the `sqlid` at the end, and although the Cursor IDs were different, the SQL Text and such like were all the same and so was the `sqlid` so it *had* to be the same SQL statement, just with numerous child cursors. Beware of this when looking in trace files!

If you know the `sqlid` then you can count and find any and all child cursors, and their IDs in the trace file with a swift `grep` command.

The following command will count the number of child cursors in the trace file for the SQL statement with `sqlid 0um91dczrf666`:

```
grep -i "0um91dczrf666" TraceFile.trc | wc -l
```

Listing 2.1: Counting Child Cursors in a Trace File

The following command will extract, in the order of appearance in the trace file, all the child cursors for the same statement as above:

```
grep -i "0um91dczrf666" TraceFile.trc | cut -c 20-39
```

Listing 2.2: Extracting Child Cursors from a Trace File

This was an HP-UX trace file, so the digits after the # in the Cursor ID were from position 20 to 39 in the output from `grep`. Other Operating Systems may differ, so test on yours first, just in case.



3. Oracle Trace Files Explained

3.1 Trace File Sections

The trace file is made up of two main sections, the header and the trace details.

3.1.1 The Header

The header is the top of the file and consists of a few lines of text giving details of where the trace file came from, which server (operating System) it was created on, various details about the server and the database and so on.

The following is an example of a trace file created on a Windows server, running Oracle 11.2.0.4. Server, database and other potentially sensitive information has been obfuscated to protect the guilty, me!

```
Trace file C:\ORACLEDATABASE\diag\rdbms\cfg\cfg\trace\  
    => orcl_ora_27680_FREESPACE.trc  
Oracle Database 11g Enterprise Edition Release 11.2.0.4.0 - 64 bit  
    => Production  
Windows NT Version V6.2  
CPU : 8 - type 8664, 8 Physical Cores  
Process Affinity : 0x0x0000000000000000  
Memory (Avail/Total): Ph:29917M/57343M, Ph+PgF:24634M/65535M  
Instance name: orcl  
Redo thread mounted by this instance: 1  
Oracle process number: 373  
Windows thread id: 27680, image: ORACLE.EXE (SHAD)
```

```
*** 2017-06-27 13:56:36.872  
*** SESSION ID:(1017.1085) 2017-06-27 13:56:36.872  
*** CLIENT ID:() 2017-06-27 13:56:36.872  
*** SERVICE NAME:(ORCLSRV) 2017-06-27 13:56:36.872
```



```
*** MODULE NAME:(TOAD background query session) 2017-06-27 13:56:36.872
*** ACTION NAME:() 2017-06-27 13:56:36.872
```

```
=====
```

Listing 3.1: Oracle 11g Trace File Header

The last line, consisting of equals signs, is the separator between the header and the following trace details.

In Oracle versions before 10g, the latter chunk of text is not found, the one detailing session, client, module etc, those first appeared at 10g.

Note

Sometimes, you might see a header that has a number of detail lines before the final separator, for example:

```
...
*** 2017-05-08 12:39:13.496
*** SESSION ID:(777.2309) 2017-05-08 12:39:13.496
*** CLIENT ID:() 2017-05-08 12:39:13.496
*** SERVICE NAME:(SYS$USERS) 2017-05-08 12:39:13.496
*** MODULE NAME:(SQL*Plus) 2017-05-08 12:39:13.496
*** ACTION NAME:() 2017-05-08 12:39:13.496

WAIT #523653448: nam='SQL*Net message to client' ela= 1
    => driver id=1111838976 #bytes=1 p3=0 obj#=14232 tim
    => =1743670562625
WAIT #523653448: nam='SQL*Net message from client' ela= 1027
    => driver id=1111838976 #bytes=1 p3=0 obj#=14232 tim
    => =1743670564245
CLOSE #523653448: c=0,e=7,dep=0,type=1,tim=1743670564280
WAIT #0: nam='SQL*Net more data from client' ela= 12 driver
    => id=1111838976 #bytes=10 p3=0 obj#=14232 tim
    => =1743670564312
=====
```

Listing 3.2: Oracle 11g Trace File Header - With Waits etc

In this case, the trace was started after the session had parsed a query, and was in the process of executing it (EXEC and, if necessary FETCH) when the trace started. Oracle has not shown the PARSING IN CURSOR line(s) for the query in question, which is a shame.

Normally, you can ignore the detail lines above the header separator, however, you sometimes see numerous WAIT lines above the separator, in which case, there could be something to investigate. This is more likely to occur when you start tracing a session that is taking a huge amount of time, and is already under way when you begin the trace.

3.1.2 Trace Details

The majority of the trace file consists of the full trace details for the session that was traced. There are numerous lines of text here, each different, each with their own fields of one kind or another.

These are explained in the following sections.

Note

Remember, when Oracle writes a line to the trace file, it is done at the *end* of the process that was indicated by the written line. For example, if you see a FETCH line in the trace file, that

was written at the time (*tim*) that the FETCH completed. The trace file should be considered a list of things *that have happened* and not *things that are about to happen*.

For example, a recent trace file that I had the pleasure of examining had over 2 million lines. The problem was to determine where a statement was losing over 2,000 seconds of time on its first EXEC which it did not have on any of the subsequent EXECs.

The statement was parsed at *dep=0* on line 36,656 but was not executed, again at *dep=0* until line 284,530, as the following *grep* output shows:

```
grep -n "dep=0" tracefile.trc
...
36656:PARSE #11529215045668893824:c=0,e=1076,p=0,cr=0,cu=0,mis=1,r=0,
    => dep=0,og=1,plh=0,tim=26375808573637
284530:EXEC #11529215045668893824:c=83870000,e=2048188431,p=1233,cr
    => =45536,cu=18,mis=1,r=1,dep=0,og=1,plh=0,tim=26377856794348
...
```

Listing 3.3: Oracle 11g Problem Trace - *grep* output

From the above you can see that nothing took place between the PARSE and the EXEC, so any other SQL executed between those lines was done in order to facilitate the EXEC at line 284,530. So that's where the time was lost, in recursive SQL statements, executed in order to carry out the required statement's EXEC.

A quick investigation showed that Oracle was gathering dynamic stats against numerous dictionary tables in the SYS schema, so after a quick *DBMS_STATS.GATHER_SCHEMA_STATS* and a *DBMS_STATS.GATHER_DICTIONARY_STATS* the lost time was restored and the problem went away.

We *know* that the time was being lost in the EXEC as the PARSE of the statement had been written to the trace file at line 36,656 so at that point, the PARSE had completed. The EXEC did not complete until line 284,530, so everything between those lines was related to the EXEC and nothing to do with the PARSE.

Also, the *elapsed time* for the EXEC shows that the execution of the statement took 2048.188431 seconds (*e=2048188431*).

Timestamp Lines

One line that you should be interested in is this one from the header above:

```
*** 2017-06-27 13:56:36.872
```

Listing 3.4: Time Stamp Line

This is the first timestamp line in the trace file and sets the baseline for all the *tim* fields (these will be explained below) that follow, however, briefly, the *tim* values are in microseconds (millionths of a second) from a specific “epoch” - which depends on the operating system - and there isn't a consistent, operating system independent, method of converting *tim* values from a huge number of microseconds to an actual date and time that humans will understand.

There are usually a few timestamp lines written to the trace file, depending on how long it has been processing for, and these mean that we can, with a bit of fiddling, relate a `tim` value to an actual time on the clock.

Tim Values

Many trace lines hold a `tim` field which holds an offset from some Operating System defined epoch. In trace files up to Oracle 9i, these values were in hundredths of a second. Since 9i, these are now in micro seconds - or millionths of a second.

The *first* `tim` value that you see, following a timestamp line, will be the microsecond equivalent to the date and time in the timestamp line.

Recursive SQL

Your SQL statements are normally executed as top-level statements, but Oracle might need to execute some (a lot!) of recursive SQL statements, in order that your statement can be processed.

If, for example, you drop a user in a database with the `drop user xxx cascade` statement, Oracle goes off and executes hundreds of separate SQL statement to find out what objects the user owns, or has privileges to, and undoes all of those before finally dropping the contents of the user and finally the user itself.

Top-level SQL statements are identified by having a depth of zero. This can be seen in many of the trace file lines as `dep=0` in the various lines of the trace file.

Recursive statements, executed in the background, have a depth greater than zero, and some of these require recursive statements of their own, and so on.

This recursion leads to a foible in the trace file, your statement appears last and all the possibly nested, recursive statements will normally appear first. This is simply because in order for your statement to be executed, the recursive statements have to run to completion *first*.

For example, in a trace file I have open in front of me, the first statement with a `dep=0` occurs at line 709 in the file. Everything prior to that runs at `dep=3`, `dep=2` or `dep=1` and complete before I can see my own SQL statement.

Under normal circumstances, a statement that is parsed (executed etc) at `dep=n`, where $n > 0$, has been called recursively, to facilitate a statement, that will follow in the trace file, that is itself parsed (executed etc) at `dep=n - 1`.

Waits

WAIT lines in a trace file are similar, in that the WAIT must complete, and so is written to the trace file, *before* the statement that incurred the wait. For example, a FETCH that had to wait for `db file scattered read` events, will appear later in the trace file than the individual WAIT lines that the FETCH suffered from.

Cursor Ids

Every time you see a '#' followed by a number, you are looking at a cursor ID. In previous versions of Oracle, these were simply an ever increasing number, starting from 1 and increasing by 1 for each new cursor.

In Oracle 11g, the cursor *appears* to be an address in memory¹, and *will be reused* as cursors are closed and new ones opened. You cannot assume, therefore, that a cursor with a specific ID at the end of the trace file, relates to any other lines with that same ID previously written to the trace file, without checking for any intervening CLOSE lines with the same ID - that's just how it is now!

Active or Inactive

You might want to look in the STATUS column in V\$SESSION to determine if the session is active - doing *something* - or not. However, you must be aware that the STATUS column only gives an active reading when the statement is in a PARSE, EXEC or FETCH phase of execution. If the statement is in a WAIT, for example, it will show as inactive - but it is actually still processing the current statement. Something to bear in mind.

3.2 Trace File Details

As mentioned above, the trace file details, which consists of many different entries, is discussed in the following chapter.

¹But don't quote me on this, I saw it written down somewhere on the Oracle Support web site, but now that I need it, I cannot find it again. Sigh!

4. Trace File Entries

4.1 PARSING IN CURSOR

This is *usually* the first line you will see for a cursor. It shows the full SQL statement between the `PARSING IN CURSOR` line and the `END OF STMT` line. The SQL is displayed exactly as the user (or application) entered it. However, if the statement is invalid, or cannot be parsed, you will not see this, or the `PARSE` line in the trace file, you will only see a `PARSE ERROR` line instead.

This is not the actual `PARSE` for the cursor though, that normally follows on later, usually!

As an example, here is the `PARSING IN CURSOR` line for the SQL query that Toad runs in the background to extract the free space used in the database by various tablespaces, including temporary ones. However, I'm not showing the SQL here:

```
PARSING IN CURSOR #3220341128 len=3081 dep=0 uid=0 oct=3 lid=0 tim
=> =3520788574727 hv=3219027813 ad='7ffcb6778350' sqlid='7
=> bwtj5azxwxv5'
```

Listing 4.1: Parsing In Cursor Line

The various fields defined, and their descriptions can be seen in the table below.

| Code | Description |
|-------|--|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| len | The size, in bytes? Characters? of the SQL statement. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| uid | The user id of the user parsing the statement. |
| oct | Oracle Command Code of the SQL Statement. (See Appendices.) |
| lid | The user id of the proxy user, that the session logged in via. Usually the same as the UID. |

Table 4.1: Parsing in Cursor - Fields...*continues on next page*

| Code | Description |
|-------|---|
| tim | Time, in microseconds, that the details were written out to the trace file. |
| hv | Hash Value for the statement. |
| ad | SQLTEXT address - see V\$SQLAREA and/or V\$SQLTEXT. |
| sqlid | The SQL ID for the statement. |

Table 4.1: Parsing in Cursor - Fields

The `uid` field references the column `USER_ID` in the views `DBA_USERS`, `ALL_USERS` and `USER_USERS`.

The `lid` field also references the column `USER_ID` in the views `DBA_USERS`, `ALL_USERS` and `USER_USERS`. However, this only has a value when when logged in via a proxy user. This is the `USER_ID` of the username in the square brackets when you login - `connect norman[dunbar]` . . . both, as far as I have seen, the `uid` and `lid` values are the same, and both refer to the 'dunbar' user, not the 'norman' user. Given the same example as above, but logged in as a proxy user, we see something like the following:

```
PARSING IN CURSOR #3220341128 len=3081 dep=0 uid=220 oct=3 lid=220 ...
```

Listing 4.2: Parsing In Cursor Line - Proxy User

Both the `uid` and the `lid` parameters are the same, 220, and this is the `USER_ID` for the 'dunbar' user, although I logged in as the 'norman' user - but that logs in as a proxy for the 'dunbar' user, as in `connect norman[dunbar]` . . . I have yet to see the `uid` and `lid` numbers differ when logged in as a proxy user - so far at least.

As mentioned, the cursor ID field has a value that may (or may not) be an address in memory. However, that's not the same as the `ad` field, which is (I think) an address in memory for the cursor.

The `sqlid` field is the same as the `SQL_ID` column in `V$SQL`. The `hv` field is the hash value that Oracle used to determine if this statement was to be found in the cache or not. It matches the `SQL_HASH_VALUE` (and `PREV_HASH_VALUE`) column in `V$SESSION` and also the `HASH_VALUE` column in `V$SQL`.

Sometime Oracle will not write these lines to the trace file. If the cursor has been parsed, and subsequently closed, then re-parsed, you will see a `PARSE` line but not a new `PARSING IN CURSOR` for the statement. This *appears* to be only on those occasions where the statement is re-parsed *immediately* after being closed. I have not seen this "feature" when the cursor ID was used by a different statement in between.

4.2 PARSE

Normally, after the `PARSING IN CURSOR` lines, you will see a `PARSE` line for the same cursor ID. This is not always the case, for example, if you started the trace after the statement had been parsed, Oracle *may* write the `PARSING IN CURSOR` lines, but not the `PARSE`, to the trace file.

Other times when the `PARSE` line will not be seen is when the cursor associated with the SQL statement was previously `CLOSEd` but Oracle decided not to hard close the cursor but instead cached it for future use. Normally this happens if `SESSION_CACHED_CURSORS` is non-zero and the SQL statement has already been `PARSEd` at least three times.

See the type parameter for the `CLOSE` trace line in Section 4.9 on Page 28 for details of how and when a cursor *might* be cached.

A typical `PARSE` line will look like this:

```
PARSE #491311368:c=0,e=452,p=0,cr=0,cu=0,mis=1,r=0,dep=0,og=1,plh
=> =1388734953,tim=97734887542
```

Listing 4.3: Parse Line

| Code | Description |
|-------|---|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| c | Elapsed CPU time. Microseconds. |
| e | Elapsed wall clock time, also in microseconds. |
| p | The number of physical reads (blocks) that were necessary in order to carry out this <code>PARSE</code> . |
| cr | The number of consistent reads (blocks) that were necessary in order to carry out this <code>PARSE</code> . |
| cu | The number of current reads (blocks) that were necessary in order to carry out this <code>PARSE</code> . |
| mis | Whether this statement was found in the cache (0) or not (1). Indicates whether or not a hard parse was required. Seeing a 1 is bad, usually. |
| r | Rows processed. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| og | Optimiser goal. 1 = <code>ALL_ROWS</code> , 2 = <code>FIRST_ROWS</code> , 3 = <code>RULE</code> , 4 = <code>CHOOSE</code> . Depending on your version of Oracle, you may not see some of the above. |
| plh | Execution plan hash value. |
| tim | Time, in microseconds, when the parse for the statement completed. Not the time it took. |

Table 4.2: Parse - Fields

The `plh` value is a value that corresponds to the column `PLAN_HASH_VALUE` in `V$SQL_PLAN`, `V$SQL_PLAN_STATISTICS_ALL` and `V$SQLSTATS`. (There may be other views where this value appears, depending on the Oracle version in use.)

Block Details

The p, cr and cu statistics are the usual ones for block reads, viz:

- Physical reads occur when Oracle must read a block, from disc, into the buffer cache.
- Consistent reads are the number of blocks of UNDO that Oracle applied to one or more blocks, in order to return the data in those blocks, to the state they were in when the “query” started. The query in this case, would be any or all recursive queries necessary to facilitate the PARSE.
- Current Reads are blocks that Oracle used without needing to roll them back to the start of the “query”.

The fields for an PARSE are the same as those for an EXEC and a FETCH - see below.

The r field is *interesting*. Surely only a FETCH would process some rows? I thought so too, and in all the trace files I’ve come across, I have yet to see (ok, yet to *notice*) a PARSE line with anything other than r=0. I wonder why Oracle have it as part of the PARSE? Maybe, because they use the same format of line for the PARSE and EXEC they decided just to use the same one, and set the row count for a PARSE to zero?¹

¹We may never know!

4.3 PARSE ERROR

If a PARSE goes wrong for some reason, the trace file will show something like the following:

```
PARSE ERROR #491311368:len=26 dep=0 uid=755 oct=3 lid=755 tim
    => =97734836031 err=1031
SELECT LIVE_DEV FROM SITE
```

Listing 4.4: Parse Error Line

There will not normally be a PARSING IN CURSOR or a PARSE shown in the trace file for this statement.

If there was a password in the SQL text that failed to parse, that will be obfuscated and the statement truncated for security reasons.

| Code | Description |
|-------|---|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| len | The length of the SQL statement in bytes, maybe characters depending on your character set. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| uid | The user id of the user parsing the statement. |
| oct | Oracle Command Code of the SQL Statement. (See Appendices.) |
| lid | The user id of the proxy user, that the session logged in via. Usually the same as the UID. See Section 4.1 on page 15 for more details about proxy logins. |
| tim | Time, in microseconds, when this parse error was detected. |
| err | The Oracle error code that caused the parse to fail. |

Table 4.3: Parse Error - Fields

In this case, there was an error ORA-01031: insufficient privileges as the user parsing the statement did not have the required privileges to see the table in question.

The uid field references the column USER_ID in the views DBA_USERS, ALL_USERS and USER_USERS.

4.4 EXEC

The EXEC phase of a statement's execution is when Oracle dives into the database to build a result set of the desired data. It need not be the complete result set though, so beware of that.

A typical EXEC will resemble the following:

```
EXEC #3220341128:c=0,e=101,p=0,cr=0,cu=0,mis=0,r=0,dep=0,og=1,plh
⇒ =2215247290,tim=3520788606189
```

Listing 4.5: Exec Line

Bear in mind, however, that the EXEC elapsed times do not necessarily bear any resemblance to the entire time it took for the user to get a response. That response time includes the PARSE, EXEC, all the FETCHes and all the WAITs that were encountered.

The fields in an EXEC are as follows:

| Code | Description |
|-------|---|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| c | Elapsed CPU time. Microseconds. |
| e | Elapsed wall clock time, also in microseconds. |
| p | The number of physical reads (blocks) that were necessary in order to carry out this PARSE. |
| cr | The number of consistent reads (blocks) that were necessary in order to carry out this PARSE. |
| cu | The number of current reads (blocks) that were necessary in order to carry out this PARSE. |
| mis | Whether this statement was found in the cache (0) or not (1). Indicates whether or not a hard parse was required. Seeing a 1 is bad, usually. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| r | Rows processed. |
| og | Optimiser goal. 1 = ALL_ROWS, 2 = FIRST_ROWS, 3 = RULE, 4 = CHOOSE. Depending on your version of Oracle, you may not see some of the above. |
| plh | Execution plan hash value. |
| tim | Time, in microseconds, at which this EXEC statement was completed. |

Table 4.4: Exec - Fields

The fields for an EXEC are the same as those for a PARSE (and FETCH - see below).

The *r* field is *interesting*. Surely only a FETCH would process some rows? Not necessarily. Normally, when EXECing a statement, the number of rows processed is indeed zero, but some PL/SQL, for example, returns a row count of 1. In my demonstration trace file I see 4 statements whose EXEC has *r*=1 and all of them are of the format:

```
1 begin
2   do_something ();
```

```
3 end ;
```

Listing 4.6: PL/SQL Exec Example

The “do_something()” call is to `DBMS_MONITOR.START_TRACE`, `DBMS_MONITOR.STOP_TRACE` and calls to `DBMS_OUTPUT.GET_LINE`. These statements have no `FETCH` calls.

You may wish to refer back to page 18 for details of the p, cr and cu block statistics in the EXEC line.

4.5 FETCH

After processing the EXEC call, and any associated WAITs, Oracle may FETCH some, or all, of the rows returned by the query. A FETCH line may resemble the following:

```
FETCH #3220341128:c=1812500,e=28665431,p=11137,cr=50774,cu=316,mis=0,r
⇒ =30,dep=0,og=1,plh=2215247290,tim=3520817271640
```

Listing 4.7: Fetch Line

And anyone paying close attention will notice that the time it took to FETCH 30 rows was a little excessive at 28.66 seconds. In case you are wondering, this is a query against DBA_FREE_SPACE on an 11.2.0.4 database, and it is taking so long as it is hitting the *never to be solved* [Bug 19125876](#) which affects 11.20.4 on AIX on Power Systems. Unfortunately, *this* trace file is from a Windows server. The bug was logged on 30th June 2014, last updated on 26th June 2017, *and is not yet fixed!*

| Code | Description |
|-------|---|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| c | Elapsed CPU time. Microseconds. |
| e | Elapsed wall clock time, also in microseconds. |
| p | The number of physical reads (blocks) that were necessary in order to carry out this FETCH. |
| cr | The number of consistent reads (blocks) that were necessary in order to carry out this FETCH. |
| cu | The number of current reads (blocks) that were necessary in order to carry out this FETCH. |
| mis | Whether this statement was found in the cache (0) or not (1). |
| r | Rows processed. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| og | Optimiser goal. 1 = ALL_ROWS, 2 = FIRST_ROWS, 3 = RULE, 4 = CHOOSE. Depending on your version of Oracle, you may not see some of the above. |
| plh | Execution plan hash value. |
| tim | Time, in microseconds, that this one fetch was completed. |

Table 4.5: fetch - Fields

The fields for a FETCH are also the same as those for a PARSE and EXEC.

The *mis* field is *interesting*. Surely only a PARSE would indicate a miss in the cache? Again, I think Oracle are using the same format of a trace line for PARSE, EXEC and FETCH and simply ignoring or defaulting certain, non-applicable, fields.

You may wish to refer back to page 18 for details of the p, cr and cu block statistics in the EXEC line.

4.6 WAIT

This is the meat of the trace file, usually. Under normal circumstances, it's usual for the various phases of execution (PARSE, EXEC and FETCH) to encounter some WAIT events. And I would say, personally², that these WAIT states are where the vast majority of the total response time is encountered.

A WAIT line in the trace file will resemble the following example, up until Oracle 9i at the latest:

```
WAIT #3220341128: nam='db file sequential read' ela= 2769 p1=37 p2
=> =12931 p3=1
```

Listing 4.8: Wait Line - Oracle 9i

| Code | Description |
|-------|--|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| nam | The name of the wait event encountered. |
| ela | Elapsed wall clock time, in microseconds. |
| p1 | Event parameter 1. |
| p2 | Event parameter 2. |
| p3 | Event parameter 3. |

Table 4.6: Wait - Fields for Oracle 9i

The p1, p2 and p3 parameters are different depending on the wait event name. You may require to look them up in V\$EVENT_NAME to understand what they refer to. For the event above, we see the following:

- p1 is the file number;
- p2 is the (starting) block number;
- p3 is the number of blocks requested.

You may wish to use the following query to determine the various parameters for the WAIT events seen in the trace file:

```
1 select name, parameter1, parameter2, parameter3
2 from v$event_name
3 where name = 'db file sequential read';
```

Listing 4.9: Extracting Event Names for Oracle 9i

Substitute the appropriate event name for your particular WAIT event of course. And, don't forget, it *is* case sensitive - Oracle are seriously inconsistent when it comes to naming events!

In Oracle versions 10g onwards, the format is much easier to understand as the generic p1, p2 and p3 have been replaced by something more meaningful, as follows for the same WAIT event:

```
WAIT #3220341128: nam='db file sequential read' ela= 1023 file#=3 block
=> #=12 blocks=1 obj#=-1 tim=3520817183625
```

Listing 4.10: Wait Line - Oracle 10g Onwards

²Beware of small sample sizes!

| Code | Description |
|--------|--|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| nam | The name of the wait event encountered. |
| ela | Elapsed wall clock time, in microseconds. |
| file# | For this wait, the file number. |
| block# | For this wait, the (starting) block number. |
| blocks | For this wait, the number of blocks requested. |
| obj# | If appropriate, the object in question. |
| tim | Time, in microseconds, that this one wait was completed. |

Table 4.7: Parse - Fields for Oracle 10g Onwards



The above is an example of one particular WAIT event. Other events will have different parameters, as appropriate for the particular event. At least, from 10g onwards, we don't have to look up V\$WAIT_NAME every time we find a new event in our trace file. Oracle 11g added its own fields too.

The obj# refers to a specific object, where one is involved. In the above example, the value is -1, so there is no particular object in this WAIT. We are waiting for a block off of the disc, which *may* be part of an index or a table, etc, but for the sake of this WAIT, Oracle considers that there is not a specific object involved.

Where an obj# is not -1, then the value refers to the OBJECT_ID in the DBA_OBJECTS view.

4.7 ERROR

If an error is detected while tracing a session, the trace file may show something like the following:

```
ERROR #275452960: err=31013 tim=1075688943194
```

Listing 4.11: Error Line

| Code | Description |
|-------|--|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| err | The Oracle error code that caused the error. |
| tim | Time, in microseconds, when this error was detected. |

Table 4.8: Error - Fields

In this example, the error relates to `ORA-31013: Invalid XPATH expression` and indeed, is what the user saw on their terminal at the time - and was the reason I traced it because the application wasn't good enough in explaining which object caused the problem!

ERROR lines relate to an EXEC or FETCH operations. PARSEs have their own PARSE ERROR lines in the trace file.

4.8 STAT

In order to see the STAT lines in your trace file, the TIMED_STATISTICS parameter for the database must be set to TRUE.

In Oracle 9i, there are no execution statistics or timings displayed in the op field. These are only present from 10g onwards.

The STAT lines show you *exactly* how Oracle went about getting the data back to you. This is because, what you are seeing is the actual EXECUTION PLAN and it is possible for this to be different for the plan displayed by the EXPLAIN PLAN FOR ... statement.



In a trace file, the execution plan shows what did happen, as opposed to the explain plan showing what was planned to happen, but may have changed as *stuff* was encountered during the execution.

A typical STAT output could resemble the following:

```
STAT #5141189408 id=1 cnt=1 pid=0 pos=1 obj=20 op='TABLE ACCESS BY
=> INDEX ROWID ICOL$ (cr=4 pr=0 pw=0 time=26 us cost=2 size=54 card
=> =2)'
STAT #5141189408 id=2 cnt=1 pid=1 pos=1 obj=42 op='INDEX RANGE SCAN
=> I_ICOL1 (cr=3 pr=0 pw=0 time=22 us cost=1 size=0 card=2)'
```

Listing 4.12: Stat Line

| Code | Description |
|-------|--|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| id | The identifier for this row of the explain plan may be referred to by the pid field, in nested STAT lines. Allows a hierarchy to be built. |
| cnt | The number of rows processed by this step in explain plan. |
| pid | The parent ID for this step on the plan. Should be zero on the id=1 line. |
| pos | The position of this step, within the parent steps in the plan. |
| obj | The object identifier. |
| op | The operation performed by this step. This will include additional statistical and timing figures from 10g onwards. Oracle 9i is sadly deficient in this matter! |

Table 4.9: Stat - Fields for Oracle 9i

Although the example STAT lines above show only a couple of lines, it is possible for more than one to be present under a single parent step. In this case, the pid would be the same, but the pos would be different.



Pos *should* be a sequentiality increasing number, showing which step was executed when in order to facilitate the parent step, however, I have seen trace files where there have been more than one row with the same pid *and* the same pos. This is a *bad thing* when trying to figure

out what happened when!

In 10g and above, the additional fields in the op field itself, are as follows:

| Code | Description |
|------|---|
| cr | The number of consistent reads. |
| pr | The number of physical reads. |
| pw | The number of physical writes. |
| time | The elapsed time in microseconds. |
| cost | The cost of this step, as determined by the Cost Based Optimiser. |
| size | From 11g onwards, an estimate of the size, in bytes, of the data returned by this step. |
| card | The number of rows processed. 11g onwards. |

Table 4.10: Stat - Fields for Oracle 10g Onwards

4.9 CLOSE

An example of a CLOSE line from a trace file is as follows:

```
CLOSE #3220452784: c=0, e=13, dep=0, type=0, tim=3520822918452
```

Listing 4.13: Close Line

| Code | Description |
|-------|---|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| c | Elapsed CPU time. Microseconds. |
| e | Elapsed wall clock time, also in microseconds. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| type | Close type. |
| tim | Time, in microseconds, that this statement was closed. This is the time it was actually closed not when the close started. |

Table 4.11: Close - Fields

This line is written when a cursor used for an SQL statement, is no longer required and has been closed. The elapsed times relate to the time it took to close the cursor.

The type field is used to determine how the cursor was closed. It takes the following values:

- 0 if the cursor was hard closed. This indicates that the cursor was not saved in the server side closed cursor cache for later reuse. This can be because:
 - The statement is a DDL statement. DDL statements are never cached.
 - SESSION_CACHED_CURSORS is set to zero so no caching is permitted.
 - The statement *could* be cached, but as it has not been executed often enough (three times minimum), then it has not yet been cached.
- 1 if the cursor has been cached, as opposed to properly closed, in an empty slot in the cache.
- 2 if the cursor was cached, but caused another cursor to be aged out as there were no free slots.
- 3 if the cursor was used from the cache, and on CLOSE, remains cached.

A cursor that was cached on CLOSEing may be reused for the same SQL statement at a later time during the session, in this case, there will not be a PARSE for the statement prior to the next BINDS or EXEC for the statement as it was not closed after the previous usage. When Oracle came to PARSE the SQL again, it was found to be in the cache, and thus, was still open from the previous usage.

A cursor ID that has been hard CLOSED may be re-used by a subsequent opening of a new cursor, which can be for a different statement, or for this one again.

4.10 XCTEND

This line in a trace file indicates the end of a transaction. Sometimes you will see this:

```
PARSING IN CURSOR #398131288 len=6 dep=1 uid=90 oct=44 lid=90 tim
=> =1484913807072 hv=255718823 ad='0' sqlid='8ggw94h7mvxd7'
COMMIT
END OF STMT
PARSE #398131288:c=0,e=4,p=0,cr=0,cu=0,mis=0,r=0,dep=1,og=0,plh=0,tim
=> =1484913807071
XCTEND rlbk=0, rd_only=1, tim=1484913807100
EXEC #398131288:c=0,e=27,p=0,cr=0,cu=0,mis=0,r=0,dep=1,og=0,plh=0,tim
=> =1484913807121
CLOSE #398131288:c=0,e=0,dep=1,type=3,tim=1484913807132
```

Listing 4.14: Commit Statement with Xctend Line

You can see here the entire PARSE, EXEC and CLOSE for the statement. Sometimes, you don't see the COMMIT (or ROLLBACK) statement being parsed, all you see is the XCTEND line in the trace file. Why? I have no idea - it's an Oracle thing!

You should notice in the above, that the transaction was completed *before* the EXEC completed. remember that the trace file shows things after they have finished. So the EXEC of the COMMIT statement had to wait for the transaction to end before it could complete.

The fields you will see in an XCTEND line are as follows:

| Code | Description |
|---------|---|
| rlbk | 0 indicates COMMIT, 1 is for ROLLBACK. |
| rd_only | Whether or not the transaction was read only (1) or read write (0) regardless of COMMIT or ROLLBACK being executed to end it. |
| tim | Time, in microseconds, at which the transaction ended. |

Table 4.12: Xctend - Fields

4.11 BINDS

To see the bind details in a trace file, you need to have enabled at least trace level 4 (for a 10046 event), or set `binds => true` for DBMS_SUPPORT and DBMS_MONITOR calls to start tracing. You should also be *using* binds in your SQL statement too of course!

There is a *lot* of information in the binds section of a trace file. For 9i, the following data are listed:

| Code | Description |
|---------|--|
| dtv | Data type code. (See Appendices.) |
| mxl | Maximum length of the bind variable value. (Private maximum length in parentheses.) |
| mal | Array length. Only for array binds. |
| scl | Scale. Only for numeric binds. |
| pre | Precision. Only for numeric binds. |
| oacflg | Special flag indicating bind options. |
| oacflg2 | Second part of oacflg. |
| size | Amount of memory to be allocated for this chunk of the bind. |
| offset | Offset into this chunk for this bind buffer. |
| bfp | Bind address. |
| bln | Bind buffer length. |
| avl | Actual value or array length. A value of zero = NULL, or a PL/SQL OUT bind - if the cursor's command is 47 for PL/SQL Execution. |
| flg | Bind status flag. |
| value | Value of the bind variable. |

Table 4.13: Binds - Fields for Oracle 9i

While for 10g upwards, we would expect to see the following:

| Code | Description |
|---------|--|
| oacdty | Data type code. (See Appendices.) |
| mxl | Maximum length, in bytes, of the bind variable value. (Used length in parentheses.) |
| mxlc | Maximum length, in characters, of the bind variable value. (Used length in parentheses.) |
| mal | Array length. Only for array binds. |
| scl | Scale. Only for numeric binds. |
| pre | Precision. Only for numeric binds. |
| oacflg | Special flag indicating bind options. |
| fl2 | Second part of oacflg. |
| frm | Unknown. |
| csi | Identifier code for the database's default or national character set. (See Appendices.) Only used in character/string binds. |
| siz | Size of memory to be allocated for this chunk. |
| off | Offset into the chunk of the bind buffer. |
| kxsbbfp | Bind address. |

Table 4.14: Binds - Fields for Oracle 10g Onwards. ...continues on next page

| Code | Description |
|-------|--|
| bln | Bind buffer length. |
| avl | Actual value or array length. A value of zero = NULL, or a PL/SQL OUT bind - if the cursor's command is 47 for PL/SQL Execution. |
| flg | Bind status flag. |
| value | Value of the bind variable, or a memory dump. This field will only be shown for any bind that has a non-NULL value. |

Table 4.14: Binds - Fields for Oracle 10g Onwards

Binds always number from left to right in an SQL statement. This means that the first bind found in the statement, will be listed in the trace file as BIND#0 regardless of its actual name or number in the statement.

If a bind, for example :3, is used more than once in the *same* SQL statement, then it will appear *once* in the binds list - as BIND#2, however, the statement itself will refer to it correctly as :3 each time it is used.

The oacdtly values are listed in the appendices for reference. There are quite a few of these, and you should note that experience has shown that the Oracle documentation doesn't always seem match up to the reality of a trace file.

Always be aware, however, that just because a bind is defined to be a VARCHAR2, oacdtly=01, for example, it doesn't mean that the column in the table it relates to is also a VARCHAR2. Some people write code that passes VARCHAR2 values to a DATE column - this negating the ability to ever use an index (unless a function based index is created) on that DATE column!

The mx1 fields shows how big the buffer assigned to this bind variable is, and how much of it has been used. This is measured in bytes. The mx1c shows, where applicable, the buffer size and used size in characters³. This will depend on the character set in use. If the character size is not appropriate, then mx1c will be zero.

4.11.1 Examples

The following is an 11g example of the binds section for the following (recursive) SQL statement:

```

1 select o.owner#, o.name, o.namespace, o.obj#, d.d_timestamp,
2       nvl(d.property,0), o.type#, o.subname, d.d_attrs
3 from   dependency$ d, obj$ o
4 where  d.p_obj#=:1
5 and    (d.p_timestamp=nvl(:2,d.p_timestamp) or d.property=2)
6 and    o.owner#=nvl(:3,o.owner#)
7 and    d.d_obj#=o.obj#
8 order by o.obj#
```

Listing 4.15: Example of Recursive SQL Statement

The binds section, extracted from the trace file, looks as follows:

```

BINDS  #741210192:
Bind#0
```

³At least, that's what I think it shows!

```

oacdtty=02 mxl=22(22) mxlc=00 mal=00 scl=00 pre=00
oacflg=00 fl2=0001 frm=00 csi=00 siz=56 off=0
kxsbbbf=2c2d4c58 bln=22 avl=04 flg=05
value=104305
Bind#1
oacdtty=12 mxl=07(07) mxlc=00 mal=00 scl=00 pre=00
oacflg=11 fl2=0001 frm=00 csi=00 siz=0 off=24
kxsbbbf=2c2d4c70 bln=07 avl=07 flg=01
value="6/26/2017 9:58:26"
Bind#2
oacdtty=02 mxl=22(22) mxlc=00 mal=00 scl=00 pre=00
oacflg=01 fl2=0001 frm=00 csi=00 siz=0 off=32
kxsbbbf=2c2d4c78 bln=22 avl=00 flg=01

```

Listing 4.16: Binds Lines

Looking specifically at the above, we can see the following:

- We have two NUMBER binds, and one DATE;
- Neither of the NUMBERS have a scale or precision;
- The last bind, BIND#2, is a NUMBER and is NULL (avl=00) and so has no value clause;
- Because these binds are of specific types, where the storage is always exactly as indicated by the mxl field, the private storage used for the value is always the same length. In other words, NUMBER data types are always 22 bytes while DATES are 7;
- There are no character sets applicable to NUMBER or DATE data types (csi=00).

Character (VARCHAR2, CHAR etc) binds take the following form:

```

Bind#1
oacdtty=01 mxl=32(04) mxlc=00 mal=00 scl=00 pre=00
oacflg=10 fl2=0001 frm=01 csi=31 siz=0 off=24
kxsbbbf=610cd550 bln=32 avl=04 flg=01
value="DUAL"

```

Listing 4.17: Bind Example - VARCHAR2 with WE8ISO8859P1 Characterset

Here we can see that:

- csi=31 is listed so checking with the character set list, in the appendices of this document, we see that this means that the bind is using the WE8ISO8859P1 character set;
- The value can be easily read from the trace file. Had this been in a different character set, ALUTF16 for example (csi=2000), then the value would most likely have been dumped in hexadecimal, as follows:

```

...
value = 0 44 0 55 0 41 0 4C

```

Listing 4.18: Bind Example - VARCHAR2 with ALUTF16 Characterset

- We can also see that although the maximum size of the buffer for this bind is 32, (mxl=32), only 4 bytes are in use (mxl=32(04)).

A REF_CURSOR bind, will resemble the following. It has an oacdtty of 102, which appears to confuse the value field a little, as it implies that it cannot handle a data type of 102. Hmmm.

```

Bind#3
oacdtty=102 mxl=04(04) mxlc=00 mal=00 scl=00 pre=00
oacflg=01 fl2=1000000 frm=00 csi=00 siz=0 off=176
kxsbbbf=c2a91548 bln=04 avl=04 flg=01

```

```
value=Unhandled datatype (102) found in kxsbindinf  
Dump of memory from 0x00000000C2A91548 to 0x00000000C2A9154C  
0C2A91540 00000000 [...]
```

Listing 4.19: Bind Example - REF_CURSOR

4.12 UNMAP

If the cursor in question used a temporary table then you will see an UNMAP in the trace file when the cursor is CLOSED and the locks on the temporary table are freed. The UNMAP covers the following stages of operation:

- Free the lock;
- Delete the state object;
- Free the temporary segment.

If you use to analyse trace files, the UNMAP data are accumulated into the EXECUTE statistics for the cursor. (See [Oracle note 39817.1](#).)

| Code | Description |
|-------|---|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| c | Elapsed CPU time. Microseconds. |
| e | Elapsed wall clock time, also in microseconds. |
| p | The number of physical reads (blocks) that were necessary in order to carry out this PARSE. |
| cr | The number of consistent reads (blocks) that were necessary in order to carry out this PARSE. |
| cu | The number of current reads (blocks) that were necessary in order to carry out this PARSE. |
| mis | Whether this statement was found in the cache (0) or not (1). Indicates whether or not a hard parse was required. Seeing a 1 is bad, usually. |
| r | Rows processed. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| og | Optimiser goal. 1 = ALL_ROWS, 2 = FIRST_ROWS, 3 = RULE, 4 = CHOOSE. Depending on your version of Oracle, you may not see some of the above. |
| plh | Execution plan hash value. |
| tim | Time, in microseconds, when the parse for the statement completed. Not the time it took. |

Table 4.15: Parse - Fields

The plh value is a value that corresponds to the column PLAN_HASH_VALUE in V\$SQL_PLAN, V\$SQL_PLAN_STATISTICS_ALL and V\$SQLSTATS. (There may be other views where this value appears, depending on the Oracle version in use.)

UNMAP trace lines, if you ever see one, are identical to those of PARSE, EXEC, FETCH as described above, and SORT UNMAP described below.

These lines in the trace file are related to cleaning up the temporary segments in any temporary tables used by the cursor.

4.13 SORT UNMAP

SORT UNMAP lines in the trace file are similar to the UNMAP lines described above. These, however, relate to the freeing up of temporary segments used for a sort operations.

| Code | Description |
|-------|---|
| #nnnn | The cursor ID. This may be reused if for future cursors if this one is closed, and another opened. |
| c | Elapsed CPU time. Microseconds. |
| e | Elapsed wall clock time, also in microseconds. |
| p | The number of physical reads (blocks) that were necessary in order to carry out this PARSE. |
| cr | The number of consistent reads (blocks) that were necessary in order to carry out this PARSE. |
| cu | The number of current reads (blocks) that were necessary in order to carry out this PARSE. |
| mis | Whether this statement was found in the cache (0) or not (1). Indicates whether or not a hard parse was required. Seeing a 1 is bad, usually. |
| r | Rows processed. |
| dep | Recursion level. 0 = Top-level, user, SQL. |
| og | Optimiser goal. 1 = ALL_ROWS, 2 = FIRST_ROWS, 3 = RULE, 4 = CHOOSE. Depending on your version of Oracle, you may not see some of the above. |
| plh | Execution plan hash value. |
| tim | Time, in microseconds, when the parse for the statement completed. Not the time it took. |

Table 4.16: Parse - Fields

The plh value is a value that corresponds to the column PLAN_HASH_VALUE in V\$SQL_PLAN, V\$SQL_PLAN_STATISTICS_ALL and V\$SQLSTATS. (There may be other views where this value appears, depending on the Oracle version in use.)

SORT UNMAP trace lines, if you ever see one, are identical to those of PARSE, EXEC, FETCH and UNMAP as described above. They are related to cleaning up the sort segments whose details can be seen in V\$SORT_USAGE. This applies to both temporary segments in memory or actual disc based segments where a large sort has overflowed memory onto disc.



Appendices

| | | |
|----------|--|-----------|
| A | Oracle Data Types | 39 |
| B | Oracle Command Codes | 43 |
| C | Oracle Characterset Codes | 47 |
| D | How this Book Evolved | 51 |
| D.1 | Why this Book? | |
| D.2 | Creating the Book | |
| E | TraceAdjust Utility | 53 |
| F | TraceMiner2 Utility | 55 |

A. Oracle Data Types

Bind Data Types

The `oacdt` parameter in a bind variables details determines the data type of that bind variable. This is not necessarily the data type of the column in a table that it may be being INSERTed or UPDATed into, or compared against in a WHERE clause.

There is a table of data types and codes in the *Internal Data Types* section of the *Data Types* chapter in the 12cR2 [Oracle Call Interface manual](#).

However, various other sources on the internet, and in books, seem to disagree with some of what the above table shows. In addition, I have come across many Oracle Trace files where a ROWID was code 11 and not code 69. Consistency? Who mentioned consistency?

In 11g (possibly 10g) onwards, there is a new view named `V$SQL_BIND_CAPTURE` which lists various details about binds variables for SQL cursors. Two columns are of interest here, `DATATYPE` and `DATATYPE_STRING` which we can use to convert a data type number into an actual data type name. In addition, we can look at the source code for the view named `DBA_TAB_COLS`, or renamed to `DBA_TAB_COLS_V$` in 12c.

| Code | Data Type |
|------|----------------------------|
| 0 | UNDEFINED |
| 1 | NVARCHAR2 or VARCHAR2 |
| 2 | NUMBER or FLOAT(precision) |
| 8 | LONG |
| 9 | NCHAR VARYING or VARCHAR |
| 11 | ROWID |
| 12 | DATE |
| 16 | DATE |
| 23 | RAW |

Table A.1: Bind Variable data Types. . . *continues on next page*

| Code | Data Type |
|------|--|
| 24 | LONG RAW |
| 25 | UB2 |
| 29 | B4 |
| 58 | ANYDATA, XMLTYPE or OPAQUE |
| 69 | ROWID |
| 96 | NCHAR or CHAR |
| 100 | BINARY_FLOAT |
| 101 | BINARY_DOUBLE |
| 102 | REF CURSOR |
| 105 | MLSLABEL |
| 106 | MLSLABEL |
| 108 | User defined Types |
| 111 | REF XMLTYPE |
| 112 | NCLOB or CLOB |
| 113 | BLOB |
| 114 | BFILE |
| 115 | CFILE |
| 121 | User defined and/or object TYPES |
| 122 | User defined and/or object TYPES, NESTED TABLE |
| 123 | User defined and/or object TYPES, VARRAY |
| 178 | TIME(scale) |
| 179 | TIME(scale) WITH TIME ZONE |
| 180 | TIMESTAMP(scale) |
| 181 | TIMESTAMP(scale) WITH TIME ZONE |
| 182 | INTERVAL YEAR(precision) TO MONTH |
| 183 | INTERVAL DAY(precision) TO SECOND(scale) |
| 208 | UROWID (Universal ROWID) |
| 231 | TIMESTAMP(scale) WITH LOCAL TIME ZONE |

Table A.1: Bind Variable data Types

Data type codes in the above table have been extracted from the *Oracle Call Interface manual*, numerous trace files and/or the views mentioned. There are, as noted, inconsistencies between the various sources which is why there are the occasional duplicate entry. Sadly, Oracle doesn't document some internal details as much as it should (or as much as *I* think it should!).



Data types 178 and 179 are both listed in 11g and 12c, but you try creating a table with the TIME(n) or TIME(n) WITH TIME ZONE and see what happens! They *appear* to be Java data types as that appears to be the only manual that lists them as valid data types.

For an example of inconsistencies, here is a ROWID with oacdt=11 as opposed to oacdt=69. The following example was taken from a trace file, created on Windows with Oracle 11.2.0.4:

```
PARSING IN CURSOR #5141169152 len=37 dep=1 uid=0 oct=3 lid=0 tim
=> =3520788504344 hv=1398610540 ad='7ffc0a95d898' sqlid='
=> grwydz59pu6mc'
select text from view$ where rowid=:1
END OF STMT
```



```
PARSE #5141169152:c=0,e=14640,p=0,cr=29,cu=0,mis=1,r=0,dep=1,og=4,plh
    => =0,tim=3520788504343
BINDS #5141169152:
Bind#0
  oacdt=11 mxl=16(16) mxlc=00 mal=00 scl=00 pre=00
  oacflg=18 fl2=0001 frm=00 csi=00 siz=16 off=0
  kxsbbbf=bff30890 bln=16 avl=16 flg=05
  value=00002787.000A.0001
```

Listing A.1: Bind Example - ROWID With oacdt=11

There's also an example of a REF_CURSOR bind on page 32 of this eBook showing the use of an oacdt=102.

B. Oracle Command Codes

Command Codes

The oct parameter in a `PARSING IN CURSOR` line in an Oracle trace file, determines the command that is being parsed in the SQL statement. Why we should need this, I have no idea, as the SQL text for the command will - obviously - show what command is being parsed. However. . .

The following (large) table outlines the various command codes and was extracted from an Oracle 12.1.0.2 database.

| Code | Command | Code | Command |
|------|-----------------|------|-----------------------------|
| 0 | UNKNOWN | 111 | DROP PUBLIC SYNONYM |
| 1 | CREATE TABLE | 112 | CREATE PUBLIC DATABASE LINK |
| 2 | INSERT | 113 | DROP PUBLIC DATABASE LINK |
| 3 | SELECT | 114 | GRANT ROLE |
| 4 | CREATE CLUSTER | 115 | REVOKE ROLE |
| 5 | ALTER CLUSTER | 116 | EXECUTE PROCEDURE |
| 6 | UPDATE | 117 | USER COMMENT |
| 7 | DELETE | 118 | ENABLE TRIGGER |
| 8 | DROP CLUSTER | 119 | DISABLE TRIGGER |
| 9 | CREATE INDEX | 120 | ENABLE ALL TRIGGERS |
| 10 | DROP INDEX | 121 | DISABLE ALL TRIGGERS |
| 11 | ALTER INDEX | 122 | NETWORK ERROR |
| 12 | DROP TABLE | 123 | EXECUTE TYPE |
| 13 | CREATE SEQUENCE | 128 | FLASHBACK |
| 14 | ALTER SEQUENCE | 129 | CREATE SESSION |
| 15 | ALTER TABLE | 130 | ALTER MINING MODEL |
| 16 | DROP SEQUENCE | 131 | SELECT MINING MODEL |
| 17 | GRANT OBJECT | 133 | CREATE MINING MODEL |

Table B.1: Oracle Command Codes. . . *continues on next page*

| Code | Command | Code | Command |
|------|----------------------|------|-----------------------------|
| 18 | REVOKE OBJECT | 134 | ALTER PUBLIC SYNONYM |
| 19 | CREATE SYNONYM | 135 | DIRECTORY EXECUTE |
| 20 | DROP SYNONYM | 136 | SQL*LOADER DIRECT PATH LOAD |
| 21 | CREATE VIEW | 137 | DATAPUMP DIRECT PATH UNLOAD |
| 22 | DROP VIEW | 138 | DATABASE STARTUP |
| 23 | VALIDATE INDEX | 139 | DATABASE SHUTDOWN |
| 24 | CREATE PROCEDURE | 140 | CREATE SQL TXLN PROFILE |
| 25 | ALTER PROCEDURE | 141 | ALTER SQL TXLN PROFILE |
| 26 | LOCK | 142 | USE SQL TXLN PROFILE |
| 27 | NO-OP | 143 | DROP SQL TXLN PROFILE |
| 28 | RENAME | 144 | CREATE MEASURE FOLDER |
| 29 | COMMENT | 145 | ALTER MEASURE FOLDER |
| 30 | AUDIT OBJECT | 146 | DROP MEASURE FOLDER |
| 31 | NOAUDIT OBJECT | 147 | CREATE CUBE BUILD PROCESS |
| 32 | CREATE DATABASE LINK | 148 | ALTER CUBE BUILD PROCESS |
| 33 | DROP DATABASE LINK | 149 | DROP CUBE BUILD PROCESS |
| 34 | CREATE DATABASE | 150 | CREATE CUBE |
| 35 | ALTER DATABASE | 151 | ALTER CUBE |
| 36 | CREATE ROLLBACK SEG | 152 | DROP CUBE |
| 37 | ALTER ROLLBACK SEG | 153 | CREATE CUBE DIMENSION |
| 38 | DROP ROLLBACK SEG | 154 | ALTER CUBE DIMENSION |
| 39 | CREATE TABLESPACE | 155 | DROP CUBE DIMENSION |
| 40 | ALTER TABLESPACE | 157 | CREATE DIRECTORY |
| 41 | DROP TABLESPACE | 158 | DROP DIRECTORY |
| 42 | ALTER SESSION | 159 | CREATE LIBRARY |
| 43 | ALTER USER | 160 | CREATE JAVA |
| 44 | COMMIT | 161 | ALTER JAVA |
| 45 | ROLLBACK | 162 | DROP JAVA |
| 46 | SAVEPOINT | 163 | CREATE OPERATOR |
| 47 | PL/SQL EXECUTE | 164 | CREATE INDEXTYPE |
| 48 | SET TRANSACTION | 165 | DROP INDEXTYPE |
| 49 | ALTER SYSTEM | 166 | ALTER INDEXTYPE |
| 50 | EXPLAIN | 167 | DROP OPERATOR |
| 51 | CREATE USER | 168 | ASSOCIATE STATISTICS |
| 52 | CREATE ROLE | 169 | DISASSOCIATE STATISTICS |
| 53 | DROP USER | 170 | CALL METHOD |
| 54 | DROP ROLE | 171 | CREATE SUMMARY |
| 55 | SET ROLE | 172 | ALTER SUMMARY |
| 56 | CREATE SCHEMA | 173 | DROP SUMMARY |
| 57 | CREATE CONTROL FILE | 174 | CREATE DIMENSION |
| 59 | CREATE TRIGGER | 175 | ALTER DIMENSION |
| 60 | ALTER TRIGGER | 176 | DROP DIMENSION |
| 61 | DROP TRIGGER | 177 | CREATE CONTEXT |
| 62 | ANALYZE TABLE | 178 | DROP CONTEXT |
| 63 | ANALYZE INDEX | 179 | ALTER OUTLINE |
| 64 | ANALYZE CLUSTER | 180 | CREATE OUTLINE |
| 65 | CREATE PROFILE | 181 | DROP OUTLINE |

Table B.1: Oracle Command Codes... *continues on next page*

| Code | Command | Code | Command |
|------|------------------------------|------|-----------------------------|
| 66 | DROP PROFILE | 182 | UPDATE INDEXES |
| 67 | ALTER PROFILE | 183 | ALTER OPERATOR |
| 68 | DROP PROCEDURE | 190 | PASSWORD CHANGE |
| 70 | ALTER RESOURCE COST | 192 | ALTER SYNONYM |
| 71 | CREATE MATERIALIZED VIEW LOG | 197 | PURGE USER_RECYCLEBIN |
| 72 | ALTER MATERIALIZED VIEW LOG | 198 | PURGE DBA_RECYCLEBIN |
| 73 | DROP MATERIALIZED VIEW LOG | 199 | PURGE TABLESPACE |
| 74 | CREATE MATERIALIZED VIEW | 200 | PURGE TABLE |
| 75 | ALTER MATERIALIZED VIEW | 201 | PURGE INDEX |
| 76 | DROP MATERIALIZED VIEW | 202 | UNDROP OBJECT |
| 77 | CREATE TYPE | 204 | FLASHBACK DATABASE |
| 78 | DROP TYPE | 205 | FLASHBACK TABLE |
| 79 | ALTER ROLE | 206 | CREATE RESTORE POINT |
| 80 | ALTER TYPE | 207 | DROP RESTORE POINT |
| 81 | CREATE TYPE BODY | 208 | PROXY AUTHENTICATION ONLY |
| 82 | ALTER TYPE BODY | 209 | DECLARE REWRITE EQUIVALENCE |
| 83 | DROP TYPE BODY | 210 | ALTER REWRITE EQUIVALENCE |
| 84 | DROP LIBRARY | 211 | DROP REWRITE EQUIVALENCE |
| 85 | TRUNCATE TABLE | 212 | CREATE EDITION |
| 86 | TRUNCATE CLUSTER | 213 | ALTER EDITION |
| 88 | ALTER VIEW | 214 | DROP EDITION |
| 91 | CREATE FUNCTION | 215 | DROP ASSEMBLY |
| 92 | ALTER FUNCTION | 216 | CREATE ASSEMBLY |
| 93 | DROP FUNCTION | 217 | ALTER ASSEMBLY |
| 94 | CREATE PACKAGE | 218 | CREATE FLASHBACK ARCHIVE |
| 95 | ALTER PACKAGE | 219 | ALTER FLASHBACK ARCHIVE |
| 96 | DROP PACKAGE | 220 | DROP FLASHBACK ARCHIVE |
| 97 | CREATE PACKAGE BODY | 225 | ALTER DATABASE LINK |
| 98 | ALTER PACKAGE BODY | 226 | CREATE PLUGGABLE DATABASE |
| 99 | DROP PACKAGE BODY | 227 | ALTER PLUGGABLE DATABASE |
| 100 | LOGON | 228 | DROP PLUGGABLE DATABASE |
| 101 | LOGOFF | 229 | CREATE AUDIT POLICY |
| 102 | LOGOFF BY CLEANUP | 230 | ALTER AUDIT POLICY |
| 103 | SESSION REC | 231 | DROP AUDIT POLICY |
| 104 | SYSTEM AUDIT | 232 | CODE-BASED GRANT |
| 105 | SYSTEM NOAUDIT | 233 | CODE-BASED REVOKE |
| 106 | AUDIT DEFAULT | 238 | ADMINISTER KEY MANAGEMENT |
| 107 | NOAUDIT DEFAULT | 239 | CREATE MATERIALIZED ZONEMAP |
| 108 | SYSTEM GRANT | 240 | ALTER MATERIALIZED ZONEMAP |
| 109 | SYSTEM REVOKE | 241 | DROP MATERIALIZED ZONEMAP |
| 110 | CREATE PUBLIC SYNONYM | 305 | ALTER PUBLIC DATABASE LINK |

Table B.1: Oracle Command Codes

The exact list of commands for your particular database version can be extracted using the following SQL command:

```
1 | select action as code ,
```



```
2 name as command  
3 from audit_actions ;
```

Listing B.1: SQL Query to List Oracle Command Codes

There are 212 different commands in Oracle 12c (12.1.0.2) while Oracle 11g (11.2.0.4) has only (!) 181.

C. Oracle Characterset Codes

Bind Charactersets

Some data types use different character sets. These are coded in the `csi` field in the bind details lines of the trace file. You can extract the list of current character set codes and names with the following query:

```
1 select nls_charset_id(value), value
2 from v$nls_valid_values
3 where isdeprecated='FALSE'
4 and parameter = 'CHARACTERSET'
5 order by nls_charset_id(value);
```

Listing C.1: SQL Query to list Character Set Codes and Names

The values that you may see here are as follows, taken from an Oracle 12.1.0.1 database where there are 222 character sets listed:

| Code | Character Set | Code | Character Set | Code | Character Set |
|------|---------------|------|------------------|------|-----------------|
| 1 | US7ASCII | 159 | CL8MACCYRILLICS | 301 | EE8EBCDIC870C |
| 2 | WE8DEC | 160 | WE8PC860 | 312 | TR8EBCDIC1026S |
| 3 | WE8HP | 161 | IS8PC861 | 314 | BLT8EBCDIC1112S |
| 4 | US8PC437 | 162 | EE8MACCES | 315 | IW8EBCDIC424S |
| 5 | WE8EBCDIC37 | 163 | EE8MACCROATIANS | 316 | EE8EBCDIC870S |
| 6 | WE8EBCDIC500 | 164 | TR8MACTURKISHS | 317 | CL8EBCDIC1025S |
| 7 | WE8EBCDIC1140 | 165 | IS8MACICELANDICS | 319 | TH8TISEBCDICS |
| 8 | WE8EBCDIC285 | 166 | EL8MACGREEKS | 320 | AR8EBCDIC420S |
| 9 | WE8EBCDIC1146 | 167 | IW8MACHEBREWS | 322 | CL8EBCDIC1025C |
| 10 | WE8PC850 | 170 | EE8MSWIN1250 | 323 | CL8EBCDIC1025R |
| 11 | D7DEC | 171 | CL8MSWIN1251 | 324 | EL8EBCDIC875R |

Table C.1: Oracle Character Set Codes... continues on next page

| Code | Character Set | Code | Character Set | Code | Character Set |
|------|----------------|------|-----------------|------|------------------|
| 12 | F7DEC | 172 | ET8MSWIN923 | 325 | CL8EBCDIC1158 |
| 13 | S7DEC | 173 | BG8MSWIN | 326 | CL8EBCDIC1158R |
| 14 | E7DEC | 174 | EL8MSWIN1253 | 327 | EL8EBCDIC423R |
| 15 | SF7ASCII | 175 | IW8MSWIN1255 | 351 | WE8MACROMAN8 |
| 16 | NDK7DEC | 176 | LT8MSWIN921 | 352 | WE8MACROMAN8S |
| 17 | I7DEC | 177 | TR8MSWIN1254 | 353 | TH8MACTHAI |
| 18 | NL7DEC | 178 | WE8MSWIN1252 | 354 | TH8MACTHAIS |
| 19 | CH7DEC | 179 | BLT8MSWIN1257 | 368 | HU8CW12 |
| 20 | YUG7ASCII | 180 | D8EBCDIC273 | 380 | EL8PC437S |
| 21 | SF7DEC | 181 | I8EBCDIC280 | 381 | EL8EBCDIC875 |
| 22 | TR7DEC | 182 | DK8EBCDIC277 | 382 | EL8PC737 |
| 23 | IW7IS960 | 183 | S8EBCDIC278 | 383 | LT8PC772 |
| 25 | IN8ISCII | 184 | EE8EBCDIC870 | 384 | LT8PC774 |
| 27 | WE8EBCDIC1148 | 185 | CL8EBCDIC1025 | 385 | EL8PC869 |
| 28 | WE8PC858 | 186 | F8EBCDIC297 | 386 | EL8PC851 |
| 31 | WE8ISO8859P1 | 187 | IW8EBCDIC1086 | 390 | CDN8PC863 |
| 32 | EE8ISO8859P2 | 188 | CL8EBCDIC1025X | 401 | HU8ABMOD |
| 33 | SE8ISO8859P3 | 189 | D8EBCDIC1141 | 500 | AR8ASMO8X |
| 34 | NEE8ISO8859P4 | 190 | N8PC865 | 554 | AR8NAFITHA711 |
| 35 | CL8ISO8859P5 | 191 | BLT8CP921 | 555 | AR8SAKHR707 |
| 36 | AR8ISO8859P6 | 192 | LV8PC1117 | 556 | AR8MUSSAD768 |
| 37 | EL8ISO8859P7 | 193 | LV8PC8LR | 557 | AR8ADOS710 |
| 38 | IW8ISO8859P8 | 194 | BLT8EBCDIC1112 | 558 | AR8ADOS720 |
| 39 | WE8ISO8859P9 | 195 | LV8RST104090 | 559 | AR8APTEC715 |
| 40 | NE8ISO8859P10 | 196 | CL8KOI8R | 560 | AR8MSWIN1256 |
| 41 | TH8TISASCII | 197 | BLT8PC775 | 561 | AR8NAFITHA721 |
| 42 | TH8TISEBCDIC | 198 | DK8EBCDIC1142 | 563 | AR8SAKHR706 |
| 43 | BN8BSCII | 199 | S8EBCDIC1143 | 565 | AR8ARABICMAC |
| 44 | VN8VN3 | 200 | I8EBCDIC1144 | 566 | AR8ARABICMACS |
| 45 | VN8MSWIN1258 | 201 | F7SIEMENS9780X | 590 | LA8ISO6937 |
| 46 | WE8ISO8859P15 | 202 | E7SIEMENS9780X | 829 | JA16VMS |
| 47 | BLT8ISO8859P13 | 203 | S7SIEMENS9780X | 830 | JA16EUC |
| 48 | CEL8ISO8859P14 | 204 | DK7SIEMENS9780X | 831 | JA16EUCYEN |
| 49 | CL8ISOIR111 | 205 | N7SIEMENS9780X | 832 | JA16SJIS |
| 50 | WE8NEXTSTEP | 206 | I7SIEMENS9780X | 833 | JA16DBCS |
| 51 | CL8KOI8U | 207 | D7SIEMENS9780X | 834 | JA16SJISYEN |
| 52 | AZ8ISO8859P9E | 208 | F8EBCDIC1147 | 835 | JA16EBCDIC930 |
| 70 | AR8EBCDICX | 210 | WE8GCOS7 | 836 | JA16MACSJIS |
| 81 | EL8DEC | 211 | EL8GCOS7 | 837 | JA16EUCTILDE |
| 82 | TR8DEC | 221 | US8BS2000 | 838 | JA16SJISTILDE |
| 90 | WE8EBCDIC37C | 222 | D8BS2000 | 840 | KO16KSC5601 |
| 91 | WE8EBCDIC500C | 223 | F8BS2000 | 842 | KO16DBCS |
| 92 | IW8EBCDIC424 | 224 | E8BS2000 | 845 | KO16KSCCS |
| 93 | TR8EBCDIC1026 | 225 | DK8BS2000 | 846 | KO16MSWIN949 |
| 94 | WE8EBCDIC871 | 226 | S8BS2000 | 850 | ZHS16CGB231280 |
| 95 | WE8EBCDIC284 | 230 | WE8BS2000E | 851 | ZHS16MACCGB23128 |
| 96 | WE8EBCDIC1047 | 231 | WE8BS2000 | 852 | ZHS16GBK |

Table C.1: Oracle Character Set Codes... continues on next page

| Code | Character Set | Code | Character Set | Code | Character Set |
|------|----------------|------|-----------------|------|---------------|
| 97 | WE8EBCDIC1140C | 232 | EE8BS2000 | 853 | ZHS16DBCS |
| 98 | WE8EBCDIC1145 | 233 | CE8BS2000 | 854 | ZHS32GB18030 |
| 99 | WE8EBCDIC1148C | 235 | CL8BS2000 | 860 | ZHT32EUC |
| 100 | WE8EBCDIC1047E | 239 | WE8BS2000L5 | 861 | ZHT32SOPS |
| 101 | WE8EBCDIC924 | 241 | WE8DG | 862 | ZHT16DBT |
| 110 | EEC8EUROASCII | 251 | WE8NCR4970 | 863 | ZHT32TRIS |
| 113 | EEC8EUROPA3 | 261 | WE8ROMAN8 | 864 | ZHT16DBCS |
| 114 | LA8PASSPORT | 262 | EE8MACCE | 865 | ZHT16BIG5 |
| 140 | BG8PC437S | 263 | EE8MACCROATIAN | 866 | ZHT16CCDC |
| 150 | EE8PC852 | 264 | TR8MACTURKISH | 867 | ZHT16MSWIN950 |
| 152 | RU8PC866 | 265 | IS8MACICELANDIC | 868 | ZHT16HKSCS |
| 153 | RU8BESTA | 266 | EL8MACGREEK | 871 | UTF8 |
| 154 | IW8PC1507 | 267 | IW8MACHEBREW | 872 | UTFE |
| 155 | RU8PC855 | 277 | US8ICL | 873 | AL32UTF8 |
| 156 | TR8PC857 | 278 | WE8ICL | 992 | ZHT16HKSCS31 |
| 158 | CL8MACCYRILLIC | 279 | WE8ISOICLUK | 2000 | AL16UTF16 |

Table C.1: Oracle Character Set Codes

If you see a character set code in the `csi` field, and you don't have the above list, you can determine the character set in use with the following query:

```
1 | SELECT NLS_CHARSET_NAME(1) FROM dual;
```

Listing C.2: SQL Query to Convert a `csi` Code to a Character Set Name

Correspondingly, you can go from a character set name to its `csi` code with the following query:

```
1 | SELECT NLS_CHARSET_ID('US7ASCII') FROM dual;
```

Listing C.3: SQL Query to Convert a Character Set Name to a `csi` Code



D. How this Book Evolved

D.1 Why this Book?

This eBook originally started as a collection of useful notes, scripts, etc which I had come across, written or deduced, over the years of playing with Oracle Trace Files and various Trace File Analysis tools.

Eventually, I collected them all together in one place so that I could have all the useful information, in one easy to find place¹.

There are many places on the internet, people I have met or worked with, books which I have bought (or had bought for me) and read, to whom I will always be grateful as they have taught me many things. Maybe I can help others as I have been helped.

As Isaac Newton is thought to have said, “*I have stood on the shoulders of giants*” - me too.

D.2 Creating the Book

The book was created in plain text files, originally in **ReStructuredText** mode, written (and edited) on both Windows², or on Linux - which I use for all my personal and business needs at home.

These RST files were simply a quick way to gather notes. They were then later enhanced by adding more detail and/or example code, and converted to \LaTeX by judicious use of the **Pandoc** text conversion utility, which I *highly* recommend.

The converted \LaTeX files were then further enhanced, indexed, tidied up, etc using **TexStudio** which runs on both Windows and Linux, so I had the same development environment in both locations. Handy.

¹I'm getting old, every day it seems. I am starting to forget things I knew - they are ageing out of the cache! I now appear to have the memory capacity of a small newt, so I need to have things written down!

²At work, in my lunch hour

The book itself, was created using a L^AT_EX template called [the] [Legrand Orange Book Template](#) created by Mathias Legrand. It is thanks to him that you get this book for free, because the licence terms of the book template specify no commercial use. I'm happy with that myself.

The front cover image on this book is taken from the book *Kunstformen der Natur* by German biologist Ernst Haeckel. The book was published between 1899 and 1904. The image used is of various *Discomedusae* which are a taxonomic group of jellyfish.

You can read about them on [Wikipedia](#) and there is a brief overview of the above book, also on [Wikipedia](#), which shows a number of other images taken from the book. (Some of which I considered before choosing the current one!)

*Discomedusae*³ have absolutely nothing to do with Oracle or Trace Files - but I liked the image, and decided that it would make a good cover for the book and a decent enough chapter heading image too.

³DiscoMedusae? Sounds like dancing jellyfish to me!



E. TraceAdjust Utility

TraceAdjust Introduction

TraceAdjust is a utility that I wrote to help me process the myriads of trace files that I come across in my DBA work. You can get the source code from github in my [TraceAdjust repository](#) and compile it on Windows or Linux/Unix with any decent C++ compiler.

It reads a normal trace file and writes out an adjusted one, as follows:

- The `tim` values are converted to seconds, by inserting a decimal point in the appropriate position;
- It adds a `delta` to each `tim` line. The `delta` is the number of microseconds between the `tim` on this line, and the `tim` on the previous (appropriate) trace line. This allows me to see how long passed between the previous `tim` and this one. Occasionally useful.
- It adds a `ds1t` to each `tim` line. This is the “delta since last timestamp” and simply counts up the number of microseconds that have passed since the trace file last produced a timestamp line similar to the one in the header. Again, occasionally useful.
- It adds a `local` to each `tim` line. This is a conversion of the `tim` value on the line, to an actual date, in the current local timezone. The time part is resolved down to microsecond level. This is usually very useful!

Running a trace file through *TraceAdjust* will create a new trace file, which some trace analysing utilities cannot cope with due to the additional fields that I have introduced. The Trace File Browser, in Toad, on the other hand, copes with my trace files quite happily and simply ignores the additional data as appropriate.

The example below shows the before and after state of a `PARSE` line from a trace file:

```
PARSE #4474286416:c=0,e=418,p=0,cr=0,cu=0,mis=1,r=0,dep=1,og=4,plh=0,  
⇒ tim=1030574627220
```

Listing E.1: TraceAdjust Example - Before Processing

Which becomes:

```
PARSE #4474286416:c=0,e=418,p=0,cr=0,cu=0,mis=1,r=0,dep=1,og=4,plh=0,  
    => tim=1030574.627220,delta=-1,dslt=768371,local='2017 Mar 13  
    => 09:23:21.768371'
```

Listing E.2: TraceAdjust Example - After Processing



F. TraceMiner2 Utility

TraceMiner2 Introduction

So, you have a trace file, chock full of statements with lots of bind variables in use. You *need* to read through it to find out which execution of any of the statements used which actual bind values. How is this easily done?

*TraceMiner2*¹ is another utility that I wrote to help me process the myriads of trace files that I come across in my DBA work.

You can get the source code from github in my [TraceMiner2 repository](#) and compile it on Windows or Linux/Unix with any decent C++ compiler.

It reads a trace file and writes out an HTML report showing various details of the SQL in the file, showing:

- The line number of the file where the SQL statement was found at (PARSING IN CURSOR);
- The line number of the file where the PARSE statement was found at;
- The line number of the file where the BINDS details were found at;
- The line number of the file where the EXEC statement was found at;
- The depth (dep) of the statement when parsed etc.

You can choose to ignore statements over any given depth, so if you only want top-level and dep=1 statements, just request `-depth=1` on the command line.

The default HTML report appears as follows:

¹There will be no more unashamed plugs in this document, I promise.

TraceMiner2

Processing Trace File: test2.trc

| EXEC Line | PARSE Line | BINDS# Line | SQL Line | DEP | SQL Text |
|--------------|---------------|----------------|-------------|-----|---|
| 44 | 27 | 28 | 25 | 1 | select obj#,type#,ctime,mtime,stime, status, dataobj#, where owner#=90 and name='PK_RECONCILE_CAPSIL_UV' and linkname is null and subname is null |
| 60 | 53 | 54 | 51 | 1 | select audit\$,options from procedure\$ where obj#=90851 |
| 76 | 69 | 70 | 67 | 1 | select owner#,name,namespace,remoteowner,linkname,p_tir (property,0),subname,type#,d_attrs from dependency\$d, p_obj#=obj#(+) order by order# |
| 98 | 91 | 92 | 89 | 1 | select order#,columns,types from access\$ where d_obj#=# |
| 125 | 108 | 109 | 106 | 1 | select /*+ index(idl_sb4\$ i_idl_sb41) */ piece#,length |

However, you can, if you have a standard report format at your company, configure the generated css file to match that of your format. *TraceMiner2* will not overwrite the css file if one exists in the output folder.